

05124265: Reinforcement Learning
Exercise 3

Tal Grossman, 201512282, Moshe Yelisevitch, 207423104

10/07/2024

1 Theory

for theory sections please see handwritten solution.

RL-Ex 3.

Q1. Yes. Proof:

$$\pi^*(s) \stackrel{\text{by definition}}{=} \operatorname{argmax}_a Q_M^*(s, a) = \operatorname{argmax}_a Q_M^*(s, a) + \underbrace{f(s)}_{\text{const.}} = \operatorname{argmax}_a Q_{M'}^*(s, a)$$

$\pi^* = \operatorname{argmax}_a Q_{M'}^*(s, a)$ is defined as the ^{w.r.t. a} optimal policy

for $Q_M^*(s, a)$, so we proved that $\pi^*(s)$ is the optimal

Policy for $Q_{M'}^*(s, a)$ also!

Q3. 1. Each iteration we run π_t , the opt. policy for iter. t .

Q 2 :

- Bellman eq for m :

$$Q_m^* = \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma \max_{a'} Q_m^*(s', a')]$$

- Bellman eq for m' :

$$Q_{m'}^* = \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma \max_{a'} Q_{m'}^*(s', a')]$$

Plug in
 a' as
 a

$$= \sum_{s'} P(s'|s, a) [R(s, a, s') + \phi(s) - \gamma \phi(s') + \gamma \max_{a'} Q_m^*(s', a')]$$

$$= \sum_{s'} P(s'|s, a) R(s, a, s') + \sum_{s'} P(s'|s, a) \phi(s) - \gamma \sum_{s'} P(s'|s, a) \phi(s') \\ + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_m^*(s', a')$$

$\phi(s)$
is
independent
of s'

$$= \sum_{s'} P(s'|s, a) R(s, a, s') + \phi(s) \sum_{s'} P(s'|s, a) - \gamma \sum_{s'} P(s'|s, a) \phi(s') \\ + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_m^*(s', a')$$

$$\sum_{s'} P(s'|s, a) = 1 \Rightarrow \sum_{s'} P(s'|s, a) R(s, a, s') + \phi(s) \cdot 1 - \gamma \sum_{s'} P(s'|s, a) \phi(s') \\ + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_m^*(s', a')$$

$$\Rightarrow Q_{m'}^*(s, a) = Q_m^*(s, a) + \underbrace{\phi(s) - \gamma \sum_{s'} P(s'|s, a) \phi(s')}_{0}$$

Since this term is a state dependent and independent of a , it does not affect on the selection of the optimal policy.

\Rightarrow The optimal policy for m and m' are identical //

Q3. 1. Each iteration we run π_t , the opt. policy for iter. t ,

until a new state-action is found. Let's assume by

contradiction that exists a pair (s', a') such that

$R'_t(s', a') = 1$ we didn't discover at iteration t . In other words, the

optimal policy π_t stuck on $(s, a) : R'_t(s, a) = 0$, so $V_t^\pi = 0$.

Given a strongly connected MDP, we know (s', a') can be visited

So a policy π' that does so will get a reward of 1 at (s', a') ,

thus π is not optimal, in contradiction!

Each iteration must end!

2. $T_{\text{iteration}} = O(|S|)$. For example, in the case that we should go over all the states that we have already visited in order to go from S_0 to (s', a')

3. every state s in the set S has a set of possible action $A(s)$, So in order to go over the entire

S , a space we need $\sum_{s \in S} |A(s)| = |S| \cdot |A|$ iterations

4. $T_{\text{exploitation}} = T_{\text{iteration}} \cdot \text{num_iterations} = O(|S|) \cdot |S| \cdot |A| =$
 $= O(|S|^2 \cdot |A|)$

Q 4:

1) - from the starting position $(b, 1)$ the legal moves are:

$(b, 1) \rightarrow (a, 3)$

$(c, 3)$

$(d, 2)$

- to return to $(b, 1)$ after 2 steps the knight must move to a position which it can return to $(b, 1)$, possible moves are:

$(a, 3) \rightarrow (b, 1)$

$(c, 3) \rightarrow (b, 1)$

$(d, 2) \rightarrow (b, 1)$

- So, initially the knight has 3 legal moves

- from each it can return to $(b, 1)$

$$\Rightarrow P = \frac{1}{3} \cdot \frac{1}{8} \cdot 3 = \frac{1}{8} //$$

2) - the markov chain is irreducible since it's possible to get from any state to any other,
for the knight it can eventually get to any square from any starting position.

- the marker chain is periodic since the knight can return to any state every "pair" steps ($k \text{ is even}$), thus the biggest divisor is 2, thus periodic with $T=2$

3) the mean recurrence time T_i to a state i is the reciprocal of the stationary distribution π_i so that state

Since the knight spends an equal amount of time on each square?

$$\pi_i = 1/64 \quad \forall i$$

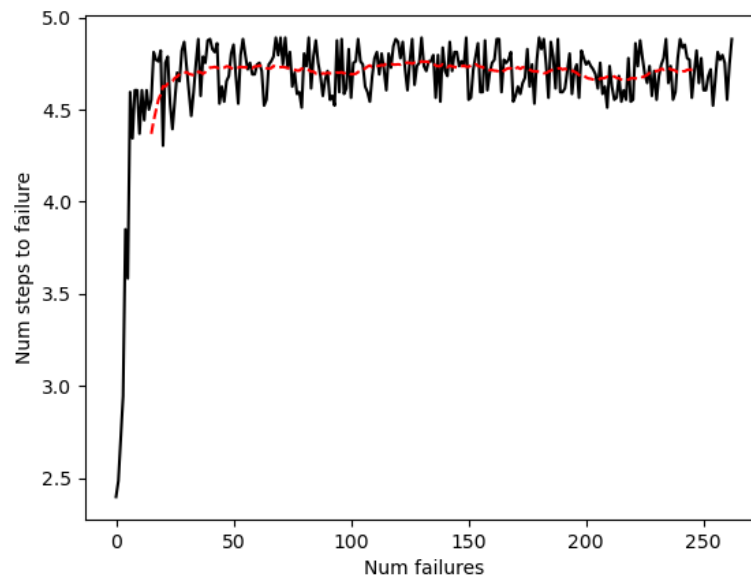
$$\Rightarrow T_{\text{corner}} = \frac{1}{\pi_{\text{corner}}} = 64$$

2 Programming

2.1 Question 1: Off-Policy Model-Based

completed in python in the attached file **control.py**

In that specific run, it took 263 iterations to converge. The plot of the failure rate is as follows:



2.2 Question 2: Q-Learning

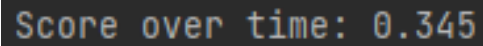
2.2.1 Item 1 - tabular_Q.py

The percentage of successful episodes is roughly 56.6% . The Q-table is as follows:

```
Score over time: 0.566
Final Q-Table Values
[[1.21427361e-01 1.83254485e-02 1.48874363e-02 1.92531811e-02]
 [1.82363151e-03 1.64267572e-03 5.90394023e-04 1.63101858e-01]
 [0.00000000e+00 3.74026090e-03 2.63652325e-03 2.57552930e-01]
 [5.99671903e-04 2.84269525e-04 2.21128911e-04 1.03078260e-01]
 [2.72628630e-01 2.31807525e-03 6.13800917e-04 1.96965210e-03]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [1.72937261e-04 8.75123858e-08 1.84143980e-01 3.24574625e-05]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [2.86188773e-03 4.22361064e-04 5.35911767e-04 3.60103396e-01]
 [0.00000000e+00 3.54639242e-01 0.00000000e+00 2.76921632e-03]
 [1.05852451e-01 1.79812038e-03 2.10456247e-04 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 2.85918529e-01 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 8.55830849e-01 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]]
```

2.2.2 Item 2 - network_Q.py

The percentage of successful episodes is roughly 34.5%. This result is worse than what we achieved with tabular Q-learning, probably due to the fact that the network is not deep and complicated enough to capture the complexity of the environment. We believe that with a deeper network with some activation functions, we could achieve better results.



```
Score over time: 0.345
```